

# An introduction to GPU computing

**José M. Mantas**

**Depto. Lenguajes y Sistemas Informáticos, Univ. Granada**

*Granada, April 15th, 2015*

**Depto. Lenguajes y Sistemas Informáticos**  
*Universidad de Granada*

# Outline

- 1 Introduction to GPU computing
- 2 CUDA Introduction
- 3 Introduction to CUDA hardware model
- 4 CUDA Programming Model
- 5 CUDA C programming Interface
- 6 Solving the 1D Linear Advection in CUDA

100

- A set of navigation icons typically found in Beamer presentations, including symbols for back, forward, search, and other slide controls.

# Graphics Processing Unit (GPU)

**Graphics Processing Unit (GPU)** Programmable single-chip processor which is used primarily for rendering of 3D graphics scenes, 3D object processing and 3D motion.




- Embedded in stand-alone cards with their **own memory** and connected to a CPU-based system (PCI-Express bus).
- **Massively parallel:** Hundreds or thousands of processing units to perform lots of arithmetic operations in parallel.
- **Highly programmable:** Non-graphics applications.
- **Excellent cost-performance ratio**

# GPUs versus multicore CPUs

Large performance gap between GPUs and multicore CPUs.

- CPUs are designed to minimize the execution latency of a single task while **GPUs are designed to maximize the total execution throughput** of a large number of parallel tasks.
- Graphics memories** are much faster.

Theoretical GFLOP/s



Source: [www.nvidia.com](http://www.nvidia.com)

# GPU computing

**GPU computing** consists of using GPUs together with CPUs to accelerate the solution of compute-intensive science, engineering and enterprise problems.

- Increasing interest in the acceleration of numerical simulations by using GPU-based computer systems.
- Initially, graphics-specific programming languages and interfaces (Cg GLSL, ...).
- NVIDIA developed the **CUDA programming toolkit** which includes an extension of the C language which facilitates the programming of GPUs.
- Additionally, **several languages and interfaces** (OpenCL, OpenACC, Thrust, ...) have been developed to make the GPU computing easier.
- **Workshop:** Acceleration of numerical algorithms using CUDA.

# Prerequisites

- Familiar with **C or C++**
- It is not necessary to know **parallel programming** but it would help.
- No knowledge about GPU architecture is needed.

# Index

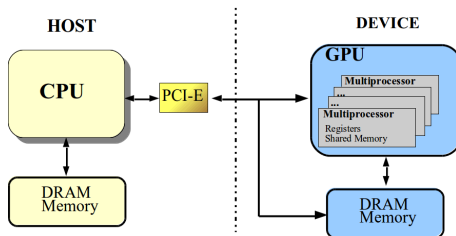
- ① Introduction to GPU computing
- ② **CUDA Introduction**
- ③ Introduction to CUDA hardware model
- ④ CUDA Programming Model
- ⑤ CUDA C programming Interface
- ⑥ Solving the 1D Linear Advection in CUDA



# CUDA (Compute Unified Device Architecture)

**CUDA:** Hardware-Software platform to exploit efficiently the potential of NVIDIA GPUs to accelerate computations.

- An **Unified architectural view** of the GPU
- A **multithreaded programming model** which includes an extension to the C language (CUDA C).
- **GPU works as a coprocessor** for the main CPU (host).
- CPU system (host) and the GPU system (device) maintain their own memories: It is possible to copy data CPU-GPU and GPU-CPU.




# Index


- ① Introduction to GPU computing
- ② CUDA Introduction
- ③ Introduction to CUDA hardware model
- ④ CUDA Programming Model
- ⑤ CUDA C programming Interface
- ⑥ Solving the 1D Linear Advection in CUDA

# Introduction to CUDA hardware model

- A CUDA enabled GPU is formed by  $N$  multiprocessors, each one having  $M$  processors or cores (32, 192, ...).
- At any clock cycle, each core of the multiprocessor executes the same instruction, but operates on different data.



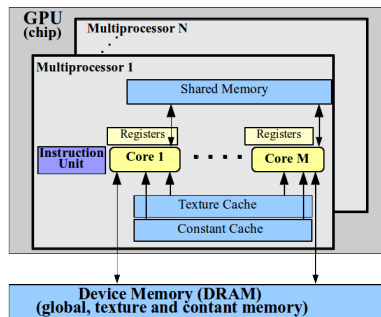
- The multiprocessors of the GPU are specialized in the parallel execution of multiple CUDA threads.
- **CUDA thread**: Sequential computational task which executes an instance of a function.
- Each CUDA thread is executed concurrently with respect to other CUDA threads (executing the same function) on the cores of the GPU multiprocessors.



# CUDA memory model

## On-chip memories:

- **Registers:** Each GPU has a particular number of registers per multiproc. which are split and assigned to the threads running on the multiproc.
- **Shared memory:** Readable and writeable only from the CUDA threads running on a multiproc. It is small but much faster than global memory.



**Off-chip memories:** Shared by all CUDA threads on the GPU.

- **Global memory:** Big but slow regarding shared memory.
- **Constant and texture memories:** read-only and cached.


# Index

- 1 Introduction to GPU computing
- 2 CUDA Introduction
- 3 Introduction to CUDA hardware model
- 4 CUDA Programming Model**
- 5 CUDA C programming Interface
- 6 Solving the 1D Linear Advection in CUDA

# CUDA kernels

## How to specify the function to be executed by each CUDA thread?


- **CUDA kernel:** special C function, called from the CPU and executed  $N$  times on the GPU by an array of  $N$  CUDA threads
- **Extremely light threads:** Recommendable to use lots of thousands of CUDA threads for high efficiency.
- Every thread executes the same code but the specific thread action depends on the **thread identifier** (built-in variable `threadIdx`).



# Simple CUDA kernel to add vectors

```
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}


...
int main()
{
    ...
    // Kernel call with N Threads
    VecAdd <<<1, N>>> (A, B, C);
}
```



- Modifier `__global__` is used to declare a kernel function.
- Each CUDA thread computes one element of the output vector `C`.
- To launch the kernel, the programmer must specify the number of CUDA threads to execute it on the GPU (parameters between `<<<` and `>>>`).



# CUDA Thread Organization. Grids and Blocks



- The set of CUDA threads for a kernel launch is organized forming a **grid of thread blocks** that run concurrently.
- When launching a kernel, the number of threads per block and the number of blocks per grid are specified with this syntax:

```
<<<  blocks_per_grid, threads_per_block >>>
```


- **Simple vector-addition CUDA program:** One grid with only one block ( $N$  threads).

# Blocks of Threads

- **CUDA block**: unit used to map threads to multiprocessors.
- Each thread is identified with a (1D, 2D or 3D) thread Index (`threadIdx`) in a **thread block**.
- Each thread block runs on a single GPU multiprocessor.
- Each multiprocessor can process a maximum number of blocks and each block can include a maximum number of threads.
- `dim3 blockDim`: stores the block dimensions for a kernel.


## 1D-Block

(`blockDim.x=4`, `blockDim.y=1`, `blockDim.z=1`)



## 2D-Block


(`blockDim.x=4`, `blockDim.y=2`, `blockDim.z=1`)



# Grids of Blocks. A general vector-addition kernel

- The blocks are organized as a (1D, 2D or 3D) matrix called **grid**.

```
#define BSIZE 256
__global__ void VecAdd(float* A, float* B, float* C, int N)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < N)    C[i] = A[i] + B[i];
}
...
int main()
{
    ...
    // Kernel call with N Threads using 256 threads 1D-blocks
    VecAdd <<<ceil((float)N/BSIZE), BSIZE>>> (A,B,C,N);
    ...
}
```



# Grids of Blocks

## How to identify parameters of a particular thread within a kernel function?

- Predefined variables

- `uint3 blockIdx`: coordinates of a block in the grid. (`blockIdx.x`, `blockIdx.y`, `blockIdx.z`).
- `dim3 blockDim`: dimensions of the block.
- `dim3 gridDim`: dimensions of the grid.

```
__global__ void VecAdd(float* A, float* B, float* C, int N)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < N)    C[i] = A[i] + B[i];
}
```

- Using these indices and `threadIdx`, programmer can specify what data are operated by each thread.
- Conditional `if` statement is needed to test if the global index values of a thread are within the valid range.

## 2D Block kernel to add matrices


- Grid and block dimensions for a kernel must be chosen depending on the particular computation to maximize efficiency.
- **Example:** 2D matrix addition, using a 2D grid of 2D-blocks.

```
...  
__global__ void MatAdd (float *A, float *B, float * C, int N)  
{  
    int j = blockIdx.x * blockDim.x + threadIdx.x; // Compute row index  
    int i = blockIdx.y * blockDim.y + threadIdx.y; // Compute column index  
    int index=i*N+j; // Compute global 1D index  
    if (i < N && j < N)  
        C[index] = A[index] + B[index]; // Computes C element if within range  
}
```

- Initially, compute the position of the current thread in the horizontal and vertical directions, using `blockIdx` and `threadIdx`.
- `i` and `j` are used to derive the global 1D index.
- 1D index is used to compute the corresponding element of `C` if `i` and `j`) are within the valid range.


# Mapping of a $45 \times 45$ matrix to a grid of $16 \times 16$ blocks

```
main()
{
    .....
    // Kernel Launch code
    dim3 threadsPerBlock (16, 16);
    dim3 numBlocks( ceil ((float)(N)/threadsPerBlock.x),
                    ceil ((float)(N)/threadsPerBlock.y) );
    MatAdd <<<numBlocks, threadsPerBlock>>> (A, B, C, N);
}
```

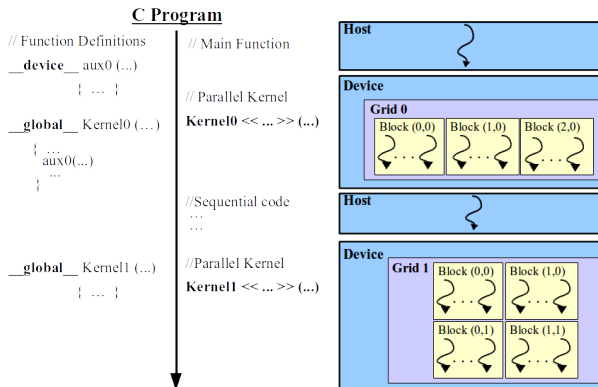


# Transparent Scalability

- Execution of several blocks can follow any execution order.
- Thread blocks are automatically mapped onto multiprocessors by the CUDA runtime system
- **Transparent scalability:** The kernel execution is automatically adapted to the number of multiprocessors of the available GPU.



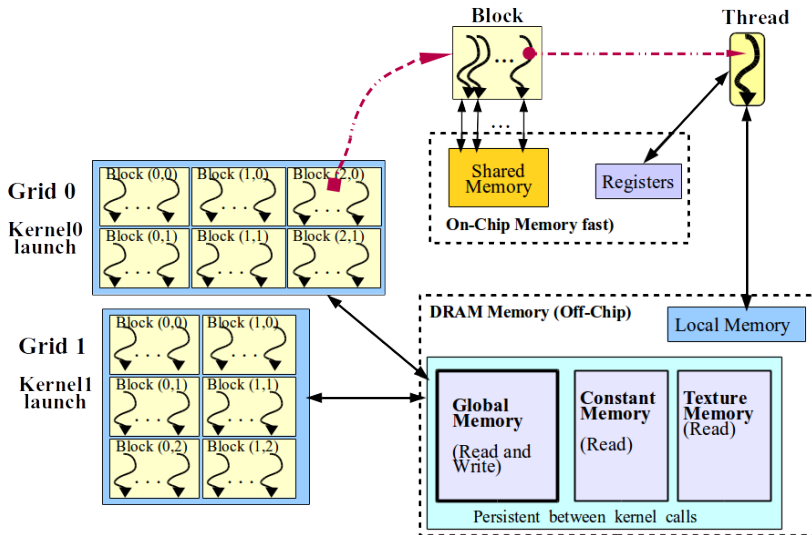
# CUDA Program Structure



- A CUDA C program must include CPU code with calls to several CUDA kernels, and the specification of these kernel functions.
- A kernel can only start on a GPU when the previous CUDA function call has returned.



# Memory Organization in a CUDA Program



# Variables in registers and local memory

## • Registers:

- All scalar variables (not arrays) declared in kernel and device functions.
- Extremely **high speed access**.
- **Scope and lifetime**: single owner thread
- **Max number of 32-bit registers** per thread (63-255) and multiprocessors (32768-65536) depends on the particular GPU.

## • Local memory:

- Array variables declared in a kernel and scalar variables which exceeding register storage capacity.
- **Same scope and lifetime** of register variables.
- Stored in global memory → **very low speed access**.

```
...
__global__ void example (float *input, float * output, int N)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x; // in register
    float w[32]; // in global memory
    ...
}
```

# Variables in Shared memory

- **Shared memory** is part of the memory space which resides on the processor chip.
- On-chip → **very High speed access**.
- **Scope**: to read and write for all the threads within the same block
- **Lifetime**: thread block activity period.
- Mean by which **threads within a block can communicate**.
- To declare a shared memory variable in a kernel or device function, we use the qualifier `__shared__`.

**Example**: Declare a shared memory vector of 32 floats in a kernel:

```
__shared__ float vector[32];
```

# Variables in Global memory

- Variables in **Global memory** can be written and read by the host by calling API functions.
- Off the processor chip, (in DRAM memory) → Speed access is much lower.
- **Lifetime:** Duration of the entire CUDA application
- Accessible from all threads through the entire CUDA application.
- The `__device__` qualifier is used to declare device variables which are resident in global memory. These variables are

```
__device__ float Global_A[32];
```

- **Main goal:** Data communication between different CUDA grids associated to different kernel calls.

# Pointers to Global memory

- We can declare pointers to global memory data in CUDA in a host function and then to allocate device memory to that pointer (with `cudaMalloc`).

**Example,** Declare and allocate a pointer to a vector of 1024 integers in global memory:


```
int numbytes = 1024*sizeof(int);  
int *a_d;  
cudaMalloc( (void**)&a_d, numbytes );
```

This pointer can be passed to a kernel function as a parameter.

**Example:** Parameters A, B and C of the `VecAdd` kernel.


# Global memory in modern GPUs

- In modern CUDA-enabled GPUs, there exists an **On-chip space memory that can be partially dedicated to cache memory** for each kernel call.
- This cache memory might reduce the global memory access cost and avoids the explicit control of the programmer.



# CUDA Thread Scheduling. Warps

- Each active block resulting from a kernel launch is divided into warps to be executed on the assigned multiprocessor.
- **Warp**: group of 32 threads with consecutive indices (ordering depends on `ThreadId`) that run physically in parallel on a multiprocessor.
- All threads in a warp execute the same instruction simultaneously.
- When a block is assigned to a multiprocessor, block is split into warps and a scheduler switches from one warp to another.



# Decomposition of a thread block into warps

- With **1D blocks**, only `threadIdx.x` is considered:

**Example:** Block size = 256 threads

First warp:  $T_0, \dots, T_{31}$ .

Second warp:  $T_{32}, \dots, T_{63}$ .

.....

Last warp (8th):  $T_{224}, \dots, T_{255}$ .

- For **2D or 3D blocks**, the dimensions are mapped to a linear 1D arrangement obtained by advancing forward firstly the  $x$ -dimension, secondly the  $y$ -dimension and finally the  $z$ -dimension.

**Example:**  $16 \times 16$  blocks

First warp:  $T_{0,0}, T_{1,0}, \dots, T_{15,0}, T_{0,1}, T_{1,1}, \dots, T_{15,1}$ .

Second warp:  $T_{0,2}, T_{1,2}, \dots, T_{15,2}, T_{0,3}, T_{1,3}, \dots, T_{15,3}$ .

.....

Last warp (8th):  $T_{0,14}, T_{1,14}, \dots, T_{15,14}, T_{0,15}, T_{1,15}, \dots, T_{15,15}$ .



# Index

- ① Introduction to GPU computing
- ② CUDA Introduction
- ③ Introduction to CUDA hardware model
- ④ CUDA Programming Model
- ⑤ **CUDA C programming Interface**
- ⑥ Solving the 1D Linear Advection in CUDA

# CUDA C programming Interface

- **CUDA C/C++** is a minimal extension of C to define kernels as C++ functions.
- To program with a high abstraction level in CUDA C: use the **CUDA runtime API** which provides functions to perform the following tasks:
  - **Manage multiple devices.**
  - **Allocate and free device memory.**
  - **Transfer data between host memory and device memory.**
- The CUDA Runtime API requires the usage of the compiler driver **nvcc**.

# CUDA Compute Capability

**CUDA Compute Capability** (CCC): indicates the architecture and features of the particular GPU.

- It is defined by a major revision number (core architecture) and a minor revision number **X.X**.

## Features of modern GPU architectures (Fermi and Kepler)

	Fermi GF100	Fermi GF104	Kepler GK104	Kepler GK110
<b>Compute Capability</b>	<b>2.0</b>	<b>2.1</b>	<b>3.0</b>	<b>3.5</b>
Threads/Warp	32	32	32	32
Max Warps/Multiproc.	48	48	64	64
Max Threads/Multiproc.	1536	1536	2048	2048
Max Blocks/Multiproc.	8	8	16	16
32-bit registers/Multiproc.	32768	32768	65536	65536
Max numb. registers/Thread	63	63	63	255
Max numb. Threads/Block	1024	1024	1024	1024
Shared memory config.	16K	16K	16K	16K
	48K	48K	48K	48K
			32K	32K

# Compiling CUDA code

- **CUDA compilation** includes **two stages**:
  - *Independent on GPU (Virtual)*: It generates lower level intermediate code (PTX) (Parallel Thread eXecution).
  - *Physical*: It generates object code for a particular GPU.
- **nvcc tool**: decouples CPU and GPU code and performs calls to the required compilers and tools: cudac, g++,...
- We can generate code for devices of an specific CUDA Compute Capability by using the -gencode compiler option:
  - **-arch=<compute\_xy>**: it generates PTX code for capability x.y.
  - **-code=<sm\\_xy>**: it generate binary code for capability x.y, by default same as -arch.

## Example:

```
nvcc filename.cu -gencode arch=compute_35, code=sm_35 -o filename
```

This command generates binary and PTX code compatible with CCC 3.5.

# Reporting errors

- All CUDA API calls return an error code (with built-in type `cudaError_t`). This error can be originated by:

- Error in the API call itself. For example:

```
cudaError_t err; err=cudaGetDevice(&devID);  
if (err!=cudaSuccess) {cout<<"ERROR!!"<<endl;}
```

- Error in an earlier asynchronous operation (for example in a kernel call):

```
kernel<<<blocksPerGrid,threadsPerBlock >>>(a,b,c);  
err = cudaGetLastError();  
if (err != cudaSuccess) {...}
```

- We can get a message which describes the particular error:

```
char *cudaGetErrorString(cudaError_t)  
printf("%s\n", cudaGetErrorString(cudaGetLastError()));
```

# Device Memory Allocation and Release

- **Allocate memory device:**  
`cudaMalloc(void ** pointer, size_t numbytes)`
- **Free memory device:** `cudaFree(void* pointer)`
- **Set a particular value in linear device memory:**  
`cudaMemset(void * pointer, int value, size_t nbytes)`

**Example:** Reset all elements of a vector with 1024 integers

```
int numbytes = 1024*sizeof(int);  
int *a_d;  
cudaMalloc( (void**)&a_d, numbytes );  
cudaMemset( a_d, 0, numbytes);  
cudaFree(a_d);
```

# Host-Device Data Transfers

**CudaMemcpy** function is used to transfer data between global memory and host memory:

```
CudaMemcpy (void *dest, void *source, size_t nbytes,  
            enum cudaMemcpyKind direction);
```

- Copies **nbytes** bytes from the memory area pointed to by **source** to the memory area pointed to by **dest**.
- The location (host or device) for **dest/source** is given by **direction**:
  - **cudaMemcpyHostToDevice**: From host to device.
  - **cudaMemcpyDeviceToHost**: From device to host.
  - **cudaMemcpyDeviceToDevice**: Between positions in device.
- The invocation blocks the CPU thread and only returns when the data copy has been completed.
- Transfer does not begin until any previous CUDA calls returned.

# Host-Device Data Transfers. Example

```
int main()
{__global__ void VecAdd(float* A, float* B, float* C,int N){ ... }
...

int N = ...;
int size = N * sizeof(float);

float* h_A = (float*) malloc(size);
float* h_B = (float*) malloc(size);
float* h_C = (float*) malloc(size);
Initialize(h_A, h_B, N);

float* d_A; float* d_B;float* d_C;
cudaMalloc((void**)&d_C, size);
cudaMalloc((void**)&d_A, size);
cudaMalloc((void**)&d_B, size);

cudaMemcpy (d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy (d_B, h_B, size, cudaMemcpyHostToDevice);

VecAdd <<<(ceil((float) N/BLOCKSIZE), BLOCKSIZE)>>> (d_A, d_B, d_C, N);

cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

cudaFree(d_A);cudaFree(d_B); cudaFree(d_C);}
```

**Host****h\_A****h\_B****h\_C****Device**



# Host-Device Data Transfers. Example

```
int main()
{__global__ void VecAdd(float* A, float* B, float* C,int N){ ... }
...
int N = ...;
int size = N * sizeof(float);

float* h_A = (float*) malloc(size);
float* h_B = (float*) malloc(size);
float* h_C = (float*) malloc(size);
Initialize(h_A, h_B, N);


float* d_A; float* d_B;float* d_C;
cudaMalloc((void**)&d_C, size);
cudaMalloc((void**)&d_A, size);
cudaMalloc((void**)&d_B, size);

cudaMemcpy (d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy (d_B, h_B, size, cudaMemcpyHostToDevice);

VecAdd <<<(ceil((float) N/BLOCKSIZE), BLOCKSIZE>>> (d_A, d_B, d_C, N);

cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

cudaFree(d_A);cudaFree(d_B); cudaFree(d_C);}
```



# Host-Device Data Transfers. Example

```
int main()
{__global__ void VecAdd(float* A, float* B, float* C,int N){ ... }
...
int N = ...;
int size = N * sizeof(float);

float* h_A = (float*) malloc(size);
float* h_B = (float*) malloc(size);
float* h_C = (float*) malloc(size);
Initialize(h_A, h_B, N);


float* d_A; float* d_B;float* d_C;
cudaMalloc((void**)&d_C, size);
cudaMalloc((void**)&d_A, size);
cudaMalloc((void**)&d_B, size);

cudaMemcpy (d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy (d_B, h_B, size, cudaMemcpyHostToDevice);

VecAdd <<<(ceil((float) N/BLOCKSIZE), BLOCKSIZE>>> (d_A, d_B, d_C, N);

cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

cudaFree(d_A);cudaFree(d_B); cudaFree(d_C);}
```



# Host-Device Data Transfers. Example

```
int main()
{__global__ void VecAdd(float* A, float* B, float* C,int N){ ... }
...
int N = ...;
int size = N * sizeof(float);

float* h_A = (float*) malloc(size);
float* h_B = (float*) malloc(size);
float* h_C = (float*) malloc(size);
Initialize(h_A, h_B, N);


float* d_A; float* d_B; float* d_C;
cudaMalloc((void**)&d_C, size);
cudaMalloc((void**)&d_A, size);
cudaMalloc((void**)&d_B, size);

cudaMemcpy (d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy (d_B, h_B, size, cudaMemcpyHostToDevice);

VecAdd <<<(ceil((float) N/BLOCKSIZE), BLOCKSIZE)>>> (d_A, d_B, d_C, N);

cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

cudaFree(d_A);cudaFree(d_B); cudaFree(d_C);}
```



# Host-Device Data Transfers. Example

```
int main()
{__global__ void VecAdd(float* A, float* B, float* C,int N){ ... }
...
int N = ...;
int size = N * sizeof(float);

float* h_A = (float*) malloc(size);
float* h_B = (float*) malloc(size);
float* h_C = (float*) malloc(size);
Initialize(h_A, h_B, N);


float* d_A; float* d_B;float* d_C;
cudaMalloc((void**)&d_C, size);
cudaMalloc((void**)&d_A, size);
cudaMalloc((void**)&d_B, size);

cudaMemcpy (d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy (d_B, h_B, size, cudaMemcpyHostToDevice);

VecAdd <<<(ceil((float) N/BLOCKSIZE), BLOCKSIZE)>>> (d_A, d_B, d_C, N);

cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

cudaFree(d_A);cudaFree(d_B); cudaFree(d_C);}
```



# Host-Device Data Transfers. Example

```
int main()
{__global__ void VecAdd(float* A, float* B, float* C,int N){ ... }
...
int N = ...;
int size = N * sizeof(float);

float* h_A = (float*) malloc(size);
float* h_B = (float*) malloc(size);
float* h_C = (float*) malloc(size);
Initialize(h_A, h_B, N);


float* d_A; float* d_B;float* d_C;
cudaMalloc((void*)&d_C, size);
cudaMalloc((void*)&d_A, size);
cudaMalloc((void*)&d_B, size);

cudaMemcpy (d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy (d_B, h_B, size, cudaMemcpyHostToDevice);

VecAdd <<<(ceil((float) N/BLOCKSIZE), BLOCKSIZE)>>> (d_A, d_B, d_C, N);

cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

cudaFree(d_A);cudaFree(d_B); cudaFree(d_C);}
```



# Index

- ① Introduction to GPU computing
- ② CUDA Introduction
- ③ Introduction to CUDA hardware model
- ④ CUDA Programming Model
- ⑤ CUDA C programming Interface
- ⑥ Solving the 1D Linear Advection in CUDA

# 1D Linear Advection

## 1D Linear Advection Equation (1DLAD)

$$\frac{\partial \phi}{\partial t} + u \frac{\partial \phi}{\partial x} = 0,$$

- $\phi = \phi(x, t)$  is a real function with  $x \in [a, b]$  and  $t > 0$ .
- Scalar  $u \in \mathbb{R}$ : propagation speed along  $x$ -direction (constant).
- Initial state:  $\phi(x, 0) = \phi_0(x)$ .
- We assume homogeneous Neumann boundary conditions

$$\frac{\partial \phi}{\partial x}(a, t) = \frac{\partial \phi}{\partial x}(b, t) = 0.$$

- Exact solution:  $\phi(x, t) = \phi_0(x - ut)$ .

# The Lax-Friedrichs method

- **Discretization** using uniform grids:

- $n + 1$  Spatial grid points:

$$x_i, \quad i = 0, \dots, n, \text{ where } x_i = a + i\Delta x, \quad \text{with } \Delta x = \frac{b - a}{n}.$$

- $M$  Time steps:  $t^1, t^2, \dots, t^M, \quad (t^{k+1} = t^k + \Delta t).$

## Lax-Friedrichs FD scheme to solve 1DLAD Eq.


$$\phi_i^{k+1} = \frac{1}{2} \left( \phi_{i-1}^k + \phi_{i+1}^k - \frac{u \Delta t}{\Delta x} (\phi_{i-1}^k - \phi_{i+1}^k) \right),$$

$$\phi_i^k = \phi(x_i, t^k), \quad i = 0, \dots, n, \quad k = 1, \dots, M.$$

- **Boundary conditions:**  $\phi_{-1}^k = \phi_0^k, \quad \phi_{n+1}^k = \phi_n^k.$



## 2-point stencil of the Lax-Friedrichs method



# A sequential C program to implement the method

- At each step, each data element of `phi_new` is computed from two neighbouring elements of `phi`.
- `phi` vector includes elements at the edges to store the ghost values.

```
void swap_pointers(float * *a, float * *b)
{float * tmp=*a;*a=*b;*b=tmp;}
...
int main(int argc, char* argv[]){
    // Define Delta x, Delta t, velocity (u) and the Courant number
    int n= ... ;float dx = ...;float dt=...; float u=...; float cu = u*dt/dx;
    //Declare the initial and final state vectors
    float * phi    =new float[n+3]; float * phi_new=new float[n+3];
    ..
    Initialize(phi); // Set phi values at t=0

    for(int k=1; k <= M; k++) // Time steps
    {// Impose Homogeneous Neumann Boundary Conditions
        phi[index(-1)] =phi[index(0)]; phi[index(n+1)]=phi[index(n)];

        for(int i=0;i<=n;i++) //Lax-Friedrichs Update
            phi_new[index(i)]=0.5*((phi[index(i+1)]+phi[index(i-1)])
                                -cu*(phi[index(i+1)]-phi[index(i-1)]));

        swap_pointers (&phi,&phi_new);
    }
```

# A CUDA Linear Advection solver (1)

- Calculation of all elements of `phi_new` for a time step can be performed in parallel.
- **i-th Thread could compute `phi_new[i]`.**
- We declare a kernel function with 4 arguments:
  - A pointer to a vector denoting the initial state vector `d_phi` ( $\phi^k$ ).
  - A pointer to the new state vector `d_phi_new` ( $\phi^{k+1}$ )
  - An integer `cu` ( $cu = \frac{u\Delta t}{\Delta x}$ )
  - An integer `n` denoting the vector size minus 2.

```
global__ void FD_kernel1 (float * d_phi,  
                           float * d_phi_new,  
                           float cu, const int n)
```

[View all posts by](#) [David J. Reardon](#)

```
#define BLOCKSIZE 256

... //FD_kernel function definition

int main(int argc, char* argv[]){
    int size=(n+3)*sizeof(float);
    float * d_phi=NULL; float * d_phi_new=NULL;
    cudaMalloc((void **) &d_phi, size); cudaMalloc((void **) &d_phi_new, size);

    Initialize(phi); // Set phi values at t=0

    cudaMemcpy(d_phi,phi,size,cudaMemcpyHostToDevice);

    for(int k=1; k <= M ;k++) // Time Steps
    {
        int blocksPerGrid =(int) ceil((float)(n+1)/BLOCKSIZE);

        // ***** CUDA Kernel Launch *****
        FD_kernel1<<<blocksPerGrid, BLOCKSIZE >>> (d_phi, d_phi_new, cu, n);
        swap_pointers (&d_phi,&d_phi_new);
    }

    cudaMemcpy(phi_GPU, d_phi, size, cudaMemcpyDeviceToHost);
}
```

# A CUDA Linear Advection solver. Kernel def.

- Each thread reads two elements of the same input vector except for two threads which read an additional element to impose the boundary conditions.
- The work of threads which are not assigned to elements of `d_phi_new` is avoided.

```
__global__ void FD_kernel1(float * d_phi,
                          float * d_phi_new,
                          float cu, int n)

int i=threadIdx.x+blockDim.x*blockIdx.x+1;

// Lax-Friedrichs Stencil
if (i<n+2)
    d_phi_new[i]=0.5*((d_phi[i+1]+d_phi[i-1])
                    -cu*(d_phi[i+1]-d_phi[i-1]));

// Impose Boundary Conditions
if (i==1)    d_phi_new[0]=d_phi_new[1];
if (i==n+1) d_phi_new[n+2]=d_phi_new[n+1];
```

# Tiling algorithms

- **Previous kernel:** The number of floating-point calculations per global memory access is low. This prevents us from achieving high performance.
- **Approach to improve:** it involves to replace global memory accesses with shared memory accesses.
- It can involve to reorganize the code to reuse data in shared memory.
- **Tiling approach** consists of two steps:
  - ① Threads of a block collaborate to load a small piece of the input data elements (*tile*), into shared memory before using them.
  - ② Threads use several times elements of the tile by accessing the shared memory.
- **Size of the tiles:** must fit into the shared memory.


# A tiled CUDA Advection solver

- All elements of `d_phi`, needed to compute all elements of `d_phi_new` for a block are loaded into shared memory.
- The number of elements to be loaded must be equal to `BLOCKSIZE + 2` (two additional elements at the edges).
- A shared memory array, `s_phi`, is declared to hold elements needed to compute `BLOCKSIZE` elements of `d_phi_new`.
- Tile is loaded in shared memory in two phases.
- After loading the tile, `__syncthreads()` function is invoked to ensure all threads in the block have finalized the load.
- When the tile is loaded for a block, we can compute the `BLOCKSIZE` elements of `d_phi_new` reading only from `s_phi`.


# Loading the tile in shared memory

Tile is loaded in shared memory in two phases.

a) **Load internal elements**



b) **Load the halo elements:** The first and last threads in the block loads the elements at the edges.





## A tiled CUDA Advection kernel (1)

```

__global__ void FD_kernel2 (float * d_phi, float * d_phi_new,
                           float cu, const int n)
{
    int li=threadIdx.x+1; //local index in shared memory vector
    int gi= blockDim.x*blockIdx.x+threadIdx.x+1; // global memory index
    int lstart=0; // start index in the block (left value)
    int lend=BLOCKSIZE+1; // end index in the block (right value)

    __shared__ float s_phi[BLOCKSIZE + 2]; //shared mem. vector
    ...
    // a) Load internal points of the tile in shared memory
    if (gi<n+2) s_phi[li] = d_phi[gi];

    // b) Load the halo points of the tile in shared memory
    if (threadIdx.x == 0) // First Thread (in the current block)
        s_phi[lstart]=d_phi[gi-1];

    if (threadIdx.x == BLOCKSIZE-1) // Last Thread
    if (gi>=n+1) // Last Block
        s_phi[(n+2)%BLOCKSIZE]=d_phi[n+2];
    else
        s_phi[lend]=d_phi[gi+1];

    __syncthreads(); // Barrier Synchronization
    ...
}

```

# A tiled CUDA Advection kernel (2)

```
__global__ void FD_kernel2 (float * d_phi, float * d_phi_new,
                           float cu, const int n)
{
    ...
    __shared__ float s_phi[BLOCKSIZE + 2]; //shared mem. vector
    float result;

    ... Load Tile (previous slide)

    __syncthreads();

    if (gi<n+2)
    {
        // Lax-Friedrichs Update
        result=0.5*((s_phi[li+1]+s_phi[li-1])
                    -cu*(s_phi[li+1]-s_phi[li-1]));
        d_phi_new[gi]=result;
    }

    // Impose Boundary Conditions
    if (gi==1) d_phi_new[0]=d_phi_new[1];
    if (gi==n+1) d_phi_new[n+2]=d_phi_new[n+1];
}
```

# Programming Resources

- **CUDA Samples:** Code samples that are included with the NVIDIA CUDA Toolkit  
<http://docs.nvidia.com/cuda/cuda-samples/>
- Sanders J., Kandrot E (2010) **CUDA by Example: An Introduction to General-Purpose GPU Programming**, Addison Wesley.
- Kirk D., Hwu Wen-mei (2012) **Programming Massively Parallel Processors, Second Edition: A Hands-on Approach**, Morgan Kaufmann.
- **CUDA C Programming Guide.** NVIDIA.  
<http://docs.nvidia.com/cuda/cuda-c-programming-guide>.
- **CUDA C Best Practices Guide.** NVIDIA.  
<http://docs.nvidia.com/cuda/cuda-c-best-practices-guide/index.html>.